



An ASGI web server, for Python.

 Test Suite passing  pypi package 0.31.0  python 3.8 | 3.9 | 3.10 | 3.11 | 3.12

Introduction

Uvicorn is an ASGI web server implementation for Python.

Until recently Python has lacked a minimal low-level server/application interface for async frameworks. The [ASGI specification](#) fills this gap, and means we're now able to start building a common set of tooling usable across all async frameworks.

Uvicorn currently supports **HTTP/1.1** and **WebSockets**.

Quickstart

Install using `pip`:

```
$ pip install uvicorn
```

This will install uvicorn with minimal (pure Python) dependencies.

```
$ pip install 'uvicorn[standard]'
```

This will install uvicorn with "Cython-based" dependencies (where possible) and other "optional extras".

In this context, "Cython-based" means the following:

- the event loop `uvloop` will be installed and used if possible.
- `uvloop` is a fast, drop-in replacement of the built-in `asyncio` event loop. It is implemented in Cython. Read more [here](#).
- The built-in `asyncio` event loop serves as an easy-to-read reference implementation and is there for easy debugging as it's pure-python based.
- the http protocol will be handled by `httptools` if possible.
- Read more about comparison with `h11` [here](#).

Moreover, "optional extras" means that:

- the websocket protocol will be handled by `websockets` (should you want to use `wsproto` you'd need to install it manually) if possible.
- the `--reload` flag in development mode will use `watchfiles`.
- windows users will have `colorama` installed for the colored logs.
- `python-dotenv` will be installed should you want to use the `--env-file` option.
- `PyYAML` will be installed to allow you to provide a `.yaml` file to `--log-config`, if desired.

Create an application:

main.py

```
async def app(scope, receive, send):
    assert scope['type'] == 'http'

    await send({
        'type': 'http.response.start',
        'status': 200,
        'headers': [
            [b'content-type', b'text/plain'],
        ],
    })
    await send({
        'type': 'http.response.body',
        'body': b'Hello, world!',
    })
```

Run the server:

```
$ uvicorn main:app
```

Usage

The uvicorn command line tool is the easiest way to run your application.

Command line options

```
$ uvicorn --help
Usage: uvicorn [OPTIONS] APP

Options:
  --host TEXT                Bind socket to this host. [default:
                             127.0.0.1]
  --port INTEGER             Bind socket to this port. If 0, an available
                             port will be picked. [default: 8000]
  --uds TEXT                 Bind to a UNIX domain socket.
  --fd INTEGER               Bind to socket from this file descriptor.
  --reload                   Enable auto-reload.
  --reload-dir PATH           Set reload directories explicitly, instead
                             of using the current working directory.
  --reload-include TEXT      Set glob patterns to include while watching
                             for files. Includes '*.py' by default; these
                             defaults can be overridden with '--reload-
                             exclude'. This option has no effect unless
                             watchfiles is installed.
  --reload-exclude TEXT      Set glob patterns to exclude while watching
                             for files. Includes '.*', .py[cod], .sw.*,
                             ~*' by default; these defaults can be
                             overridden with '--reload-include'. This
                             option has no effect unless watchfiles is
                             installed.
  --reload-delay FLOAT       Delay between previous and next check if
                             application needs to be. Defaults to 0.25s.
                             [default: 0.25]
  --workers INTEGER          Number of worker processes. Defaults to the
                             $WEB_CONCURRENCY environment variable if
                             available, or 1. Not valid with --reload.
  --loop [auto|asyncio|uvloop] Event loop implementation. [default: auto]
  --http [auto|h11|httptools] HTTP protocol implementation. [default:
                             auto]
  --ws [auto|none|websockets|wsproto] WebSocket protocol implementation.
                             [default: auto]
  --ws-max-size INTEGER       WebSocket max size message in bytes
                             [default: 16777216]
  --ws-max-queue INTEGER      The maximum length of the WebSocket message
                             queue. [default: 32]
  --ws-ping-interval FLOAT    WebSocket ping interval in seconds.
                             [default: 20.0]
  --ws-ping-timeout FLOAT     WebSocket ping timeout in seconds.
```

```

                                [default: 20.0]
--ws-per-message-deflate BOOLEAN      WebSocket per-message-deflate compression
                                        [default: True]
--lifespan [auto|on|off]              Lifespan implementation. [default: auto]
--interface [auto|asgi3|asgi2|wsgi]   Select ASGI3, ASGI2, or WSGI as the
                                        application interface. [default: auto]
--env-file PATH                       Environment configuration file.
--log-config PATH                     Logging configuration file. Supported
                                        formats: .ini, .json, .yaml.
--log-level [critical|error|warning|info|debug|trace]
                                        Log level. [default: info]
--access-log / --no-access-log        Enable/Disable access log.
--use-colors / --no-use-colors        Enable/Disable colorized logging.
--proxy-headers / --no-proxy-headers  Enable/Disable X-Forwarded-Proto,
                                        X-Forwarded-For, X-Forwarded-Port to
                                        populate remote address info.
--server-header / --no-server-header  Enable/Disable default Server header.
--date-header / --no-date-header       Enable/Disable default Date header.
--forwarded-allow-ips TEXT             Comma separated list of IP Addresses, IP
                                        Networks, or literals (e.g. UNIX Socket
                                        path) to trust with proxy headers. Defaults
                                        to the $FORWARDED_ALLOW_IPS environment
                                        variable if available, or '127.0.0.1'. The
                                        literal '*' means trust everything.
--root-path TEXT                      Set the ASGI 'root_path' for applications
                                        submounted below a given URL path.
--limit-concurrency INTEGER            Maximum number of concurrent connections or
                                        tasks to allow, before issuing HTTP 503
                                        responses.
--backlog INTEGER                     Maximum number of connections to hold in
                                        backlog
--limit-max-requests INTEGER           Maximum number of requests to service before
                                        terminating the process.
--timeout-keep-alive INTEGER           Close Keep-Alive connections if no new data
                                        is received within this timeout. [default:
                                        5]
--timeout-graceful-shutdown INTEGER    Maximum number of seconds to wait for
                                        graceful shutdown.
--ssl-keyfile TEXT                     SSL key file
--ssl-certfile TEXT                     SSL certificate file
--ssl-keyfile-password TEXT            SSL keyfile password
--ssl-version INTEGER                  SSL version to use (see stdlib ssl module's)
                                        [default: 17]
--ssl-cert-reqs INTEGER                Whether client certificate is required (see
                                        stdlib ssl module's) [default: 0]
--ssl-ca-certs TEXT                     CA certificates file
--ssl-ciphers TEXT                      Ciphers to use (see stdlib ssl module's)
                                        [default: TLSv1]
--header TEXT                          Specify custom default HTTP response headers
                                        as a Name:Value pair
--version                              Display the uvicorn version and exit.

```

<code>--app-dir TEXT</code>	Look for APP in the specified directory, by adding this to the PYTHONPATH. Defaults to the current working directory.
<code>--h11-max-incomplete-event-size INTEGER</code>	For h11, the maximum number of bytes to buffer of an incomplete event.
<code>--factory</code>	Treat APP as an application factory, i.e. a <code>() -> <ASGI app></code> callable.
<code>--help</code>	Show this message and exit.

For more information, see the [settings documentation](#).

Running programmatically

There are several ways to run uvicorn directly from your application.

`uvicorn.run`

If you're looking for a programmatic equivalent of the `uvicorn` command line interface, use `uvicorn.run()`:

`main.py`

```
import uvicorn

async def app(scope, receive, send):
    ...

if __name__ == "__main__":
    uvicorn.run("main:app", port=5000, log_level="info")
```

`Config` and `Server` instances

For more control over configuration and server lifecycle, use `uvicorn.Config` and `uvicorn.Server`:

`main.py`

```
import uvicorn

async def app(scope, receive, send):
    ...

if __name__ == "__main__":
    config = uvicorn.Config("main:app", port=5000, log_level="info")
    server = uvicorn.Server(config)
    server.run()
```

If you'd like to run Uvicorn from an already running async environment, use `uvicorn.Server.serve()` instead:

main.py

```
import asyncio
import uvicorn

async def app(scope, receive, send):
    ...

async def main():
    config = uvicorn.Config("main:app", port=5000, log_level="info")
    server = uvicorn.Server(config)
    await server.serve()

if __name__ == "__main__":
    asyncio.run(main())
```

Running with Gunicorn

**Warning**

The `uvicorn.workers` module is deprecated and will be removed in a future release.

You should use the `uvicorn-worker` package instead.

```
python -m pip install uvicorn-worker
```

Gunicorn is a mature, fully featured server and process manager.

Uvicorn includes a Gunicorn worker class allowing you to run ASGI applications, with all of Uvicorn's performance benefits, while also giving you Gunicorn's fully-featured process management.

This allows you to increase or decrease the number of worker processes on the fly, restart worker processes gracefully, or perform server upgrades without downtime.

For production deployments we recommend using gunicorn with the uvicorn worker class.

```
gunicorn example:app -w 4 -k uvicorn.workers.UvicornWorker
```

For a **PyPy** compatible configuration use `uvicorn.workers.UvicornH11Worker`.

For more information, see the [deployment documentation](#).

Application factories

The `--factory` flag allows loading the application from a factory function, rather than an application instance directly. The factory will be called with no arguments and should return an ASGI application.

main.py

```
def create_app():  
    app = ...  
    return app
```

```
$ uvicorn --factory main:create_app
```

The ASGI interface

Uvicorn uses the [ASGI specification](#) for interacting with an application.

The application should expose an async callable which takes three arguments:

- `scope` - A dictionary containing information about the incoming connection.
- `receive` - A channel on which to receive incoming messages from the server.
- `send` - A channel on which to send outgoing messages to the server.

Two common patterns you might use are either function-based applications:

```
async def app(scope, receive, send):  
    assert scope['type'] == 'http'  
    ...
```

Or instance-based applications:

```
class App:  
    async def __call__(self, scope, receive, send):  
        assert scope['type'] == 'http'  
        ...  
  
app = App()
```

It's good practice for applications to raise an exception on scope types that they do not handle.

The content of the `scope` argument, and the messages expected by `receive` and `send` depend on the protocol being used.

The format for HTTP messages is described in the [ASGI HTTP Message format](#).

HTTP Scope

An incoming HTTP request might have a connection scope like this:

```
{
    'type': 'http',
    'scheme': 'http',
    'root_path': '',
    'server': ('127.0.0.1', 8000),
    'http_version': '1.1',
    'method': 'GET',
    'path': '/',
    'headers': [
        (b'host', b'127.0.0.1:8000'),
        (b'user-agent', b'curl/7.51.0'),
        (b'accept', b'*/*')
    ]
}
```

HTTP Messages

The instance coroutine communicates back to the server by sending messages to the `send` coroutine.

```
await send({
    'type': 'http.response.start',
    'status': 200,
    'headers': [
        [b'content-type', b'text/plain'],
    ]
})
await send({
    'type': 'http.response.body',
    'body': b'Hello, world!',
})
```

Requests & responses

Here's an example that displays the method and path used in the incoming request:

```
async def app(scope, receive, send):
    """
    Echo the method and path back in an HTTP response.
    """
    assert scope['type'] == 'http'

    body = f'Received {scope["method"]} request to {scope["path"]}'
    await send({
        'type': 'http.response.start',
        'status': 200,
        'headers': [
            [b'content-type', b'text/plain'],
        ]
    })
```



```
    })
    await send({
        'type': 'http.response.body',
        'body': body.encode('utf-8'),
    })
}
```

Reading the request body

You can stream the request body without blocking the asyncio task pool, by fetching messages from the `receive` coroutine.

```
async def read_body(receive):
    """
    Read and return the entire body from an incoming ASGI message.
    """
    body = b''
    more_body = True

    while more_body:
        message = await receive()
        body += message.get('body', b'')
        more_body = message.get('more_body', False)

    return body

async def app(scope, receive, send):
    """
    Echo the request body back in an HTTP response.
    """
    body = await read_body(receive)
    await send({
        'type': 'http.response.start',
        'status': 200,
        'headers': [
            (b'content-type', b'text/plain'),
            (b'content-length', str(len(body)).encode())
        ]
    })
    await send({
        'type': 'http.response.body',
        'body': body,
    })
}
```

Streaming responses

You can stream responses by sending multiple `http.response.body` messages to the `send` coroutine.

```
import asyncio
```

```
async def app(scope, receive, send):  
    """  
    Send a slowly streaming HTTP response back to the client.  
    """  
    await send({  
        'type': 'http.response.start',  
        'status': 200,  
        'headers': [  
            [b'content-type', b'text/plain'],  
        ]  
    })  
    for chunk in [b'Hello', b' ', b'world!']:  
        await send({  
            'type': 'http.response.body',  
            'body': chunk,  
            'more_body': True  
        })  
        await asyncio.sleep(1)  
    await send({  
        'type': 'http.response.body',  
        'body': b'',  
    })
```

Why ASGI?

Most well established Python Web frameworks started out as WSGI-based frameworks.

WSGI applications are a single, synchronous callable that takes a request and returns a response. This doesn't allow for long-lived connections, like you get with long-poll HTTP or WebSocket connections, which WSGI doesn't support well.

Having an async concurrency model also allows for options such as lightweight background tasks, and can be less of a limiting factor for endpoints that have long periods being blocked on network I/O such as dealing with slow HTTP requests.

Alternative ASGI servers

A strength of the ASGI protocol is that it decouples the server implementation from the application framework. This allows for an ecosystem of interoperating webservers and application frameworks.

Daphne

The first ASGI server implementation, originally developed to power Django Channels, is [the Daphne webserver](#).

It is run widely in production, and supports HTTP/1.1, HTTP/2, and WebSockets.

Any of the example applications given here can equally well be run using `daphne` instead.

```
$ pip install daphne
$ daphne app:App
```

Hypercorn

[Hypercorn](#) was initially part of the Quart web framework, before being separated out into a standalone ASGI server.

Hypercorn supports HTTP/1.1, HTTP/2, HTTP/3 and WebSockets.

```
$ pip install hypercorn
$ hypercorn app:App
```

ASGI frameworks

You can use Uvicorn, Daphne, or Hypercorn to run any ASGI framework.

For small services you can also write ASGI applications directly.

Starlette

[Starlette](#) is a lightweight ASGI framework/toolkit.

It is ideal for building high performance asyncio services, and supports both HTTP and WebSockets.

Django Channels

The ASGI specification was originally designed for use with [Django Channels](#).

Channels is a little different to other ASGI frameworks in that it provides an asynchronous frontend onto a threaded-framework backend. It allows Django to support WebSockets, background tasks, and long-running connections, with application code still running in a standard threaded context.

Quart

[Quart](#) is a Flask-like ASGI web framework.

FastAPI

FastAPI is an API framework based on **Starlette** and **Pydantic**, heavily inspired by previous server versions of **APIStar**.

You write your API function parameters with Python 3.6+ type declarations and get automatic data conversion, data validation, OpenAPI schemas (with JSON Schemas) and interactive API documentation UIs.

BlackSheep

BlackSheep is a web framework based on ASGI, inspired by Flask and ASP.NET Core.

Its most distinctive features are built-in support for dependency injection, automatic binding of parameters by request handler's type annotations, and automatic generation of OpenAPI documentation and Swagger UI.

Falcon

Falcon is a minimalist REST and app backend framework for Python, with a focus on reliability, correctness, and performance at scale.

Muffin

Muffin is a fast, lightweight and asynchronous ASGI web-framework for Python 3.

Litestar

Litestar is a powerful, lightweight and flexible ASGI framework.

It includes everything that's needed to build modern APIs - from data serialization and validation to websockets, ORM integration, session management, authentication and more.

Panther

Panther is a fast & friendly web framework for building async APIs with Python 3.10+.

It has built-in Document-oriented Database, Caching System, Authentication and Permission Classes, Visual API Monitoring and also supports Websocket, Throttling, Middlewares.